

# Travel Product Classification — Pipeline Report

## Executive Summary

In the submitted model, I built a Keras neural network to predict whether a customer will purchase a travel product (**ProdTaken**) from a dataset of 3,208 customer records.

I first cleaned the data — fixing a gender typo (“Fe Male” → “Female”, 121 rows) and replacing outliers in **DurationOfPitch** (127 min) and **NumberOfTrips** (20–22) with median values.

I also addressed class imbalance. The original data was 81% negative / 19% positive, so I applied random oversampling to balance the classes to 50/50 (5,178 total rows). I considered using SMOTE but found simple oversampling sufficient for this dataset size.

I initially trained the baseline architecture from class (32→8→4→1) which achieved 88% accuracy, but I wanted better results. I experimented with wider networks, BatchNormalization, and learning rate scheduling.

I found that a wider, shallower network (256→128→1) with BatchNormalization and ReduceLROnPlateau significantly outperformed the baseline.

**The final model achieved 96.45% accuracy and 0.99 AUC on the held-out test set — an improvement of ~8 percentage points over the baseline.**

# Pipeline

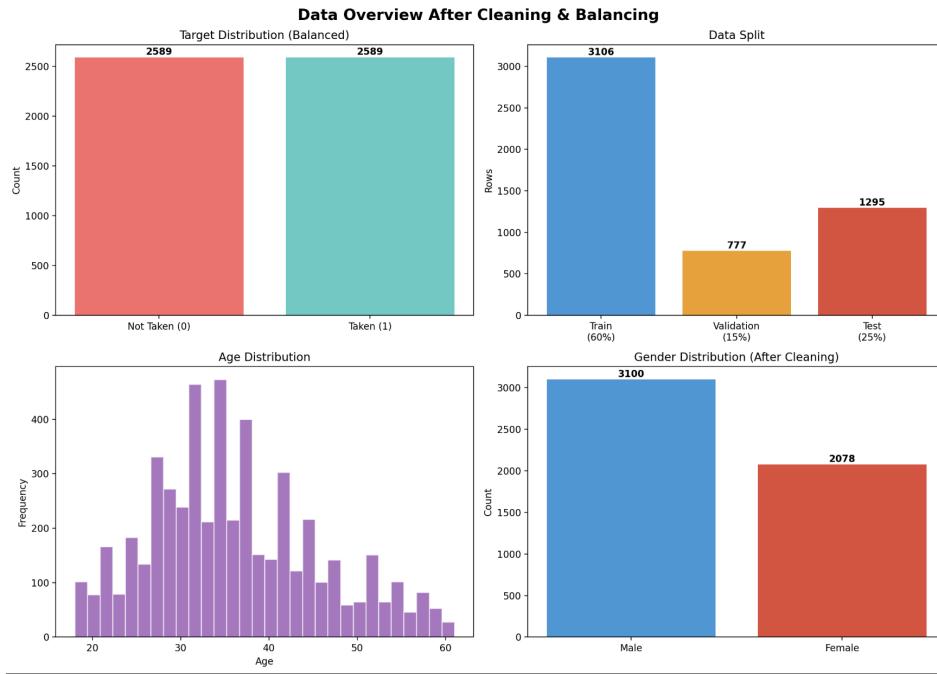
## 1. Data Cleaning & Preprocessing

The raw dataset had 3,208 rows and 19 columns. I ran exploratory data analysis and found 3 issues:

- **Gender typo:** 121 rows had “Fe Male” instead of “Female” — merged into “Female”
- **DurationOfPitch outlier:** 1 row had 127 min (next highest: 36) — replaced with median (14.0)
- **NumberOfTrips outlier:** 3 rows had values 20, 21, 22 (next highest: 8) — replaced with median (3.0)

After cleaning, I one-hot encoded the 6 categorical columns with `drop_first=True` (28 features total) and standardized all features with StandardScaler. The data was split into 3 separate files before training:

Split	Row s	%
Training	3,106	60%
Validation	777	15%
Testing	1,295	25%



To conclude: data was cleaned of typos and outliers, encoded into 28 features, and split into 3 separate files to prevent data leakage.

## 2. Class Balancing

The original target was severely imbalanced:

- ProdTaken = 0: 2,589 rows (80.7%)
- ProdTaken = 1: 619 rows (19.3%)

A naive model predicting “Not Taken” every time would be 81% accurate, making the raw data unsuitable for training.

**Solution:** Random oversampling — minority class resampled with replacement to match the majority (2,589 each → 5,178 total). The dataset was shuffled to prevent ordering bias.

To conclude: oversampling from 619 to 2,589 positive samples gave the model equal exposure to both classes and made accuracy a meaningful metric.

### 3. Model Training & Improvement

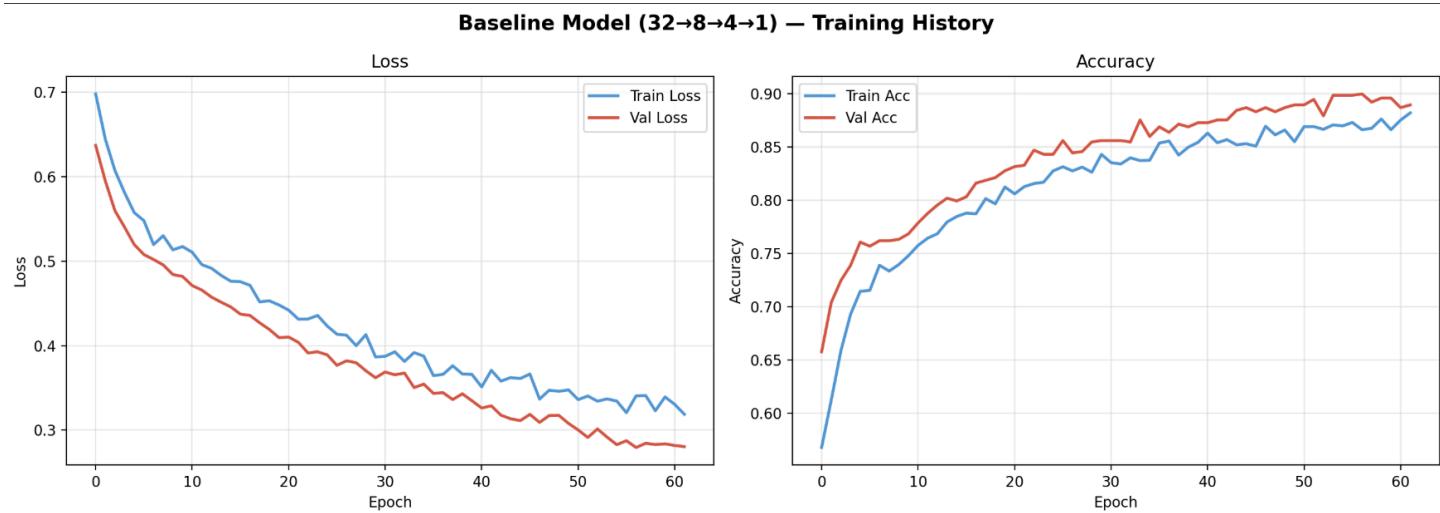
#### Baseline Model

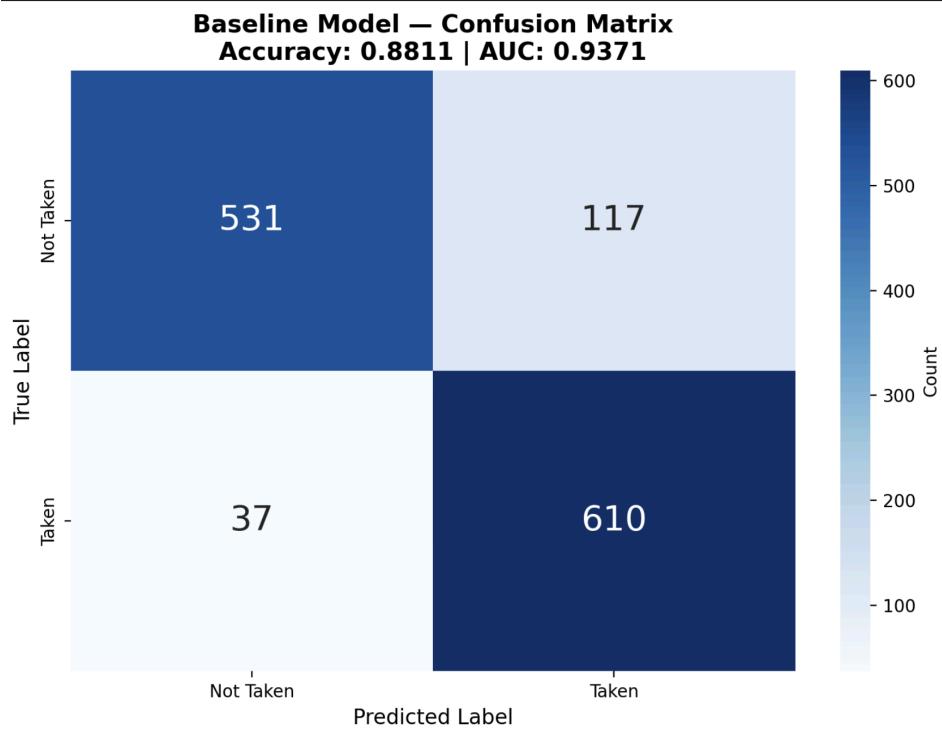
Architecture:

Dense(32, relu) → Dropout(0.3) → Dense(8, tanh) → Dense(4, relu) → Dropout(0.3) → Dense(1, sigmoid)

Trained with Adam (lr=0.001), batch\_size=16, early stopping patience=5. Ran for 62 epochs.

Result: **88.11% accuracy, 0.937 AUC**





The baseline performed okay but the validation curves suggested underfitting — the architecture didn't have enough capacity. I wanted better accuracy, so I made the following improvements:

## What I Did to Improve Accuracy

I ran 4 experiments to find the best architecture:

Experiment	Architecture	Accuracy
Baseline (class)	32→8→4→1	88.11%
Exp 1: Wider	64→32→16→1 + BatchNorm	94.21%
Exp 2: Deep + ReduceLR	128→64→32→16→1 + BatchNorm	95.06%

Exp 3: LeakyReLU variant    128→64→32→16→1 + LeakyReLU    93.59%

**Exp 4: Wide shallow**    **256→128→1 + BatchNorm**    **96.14%**

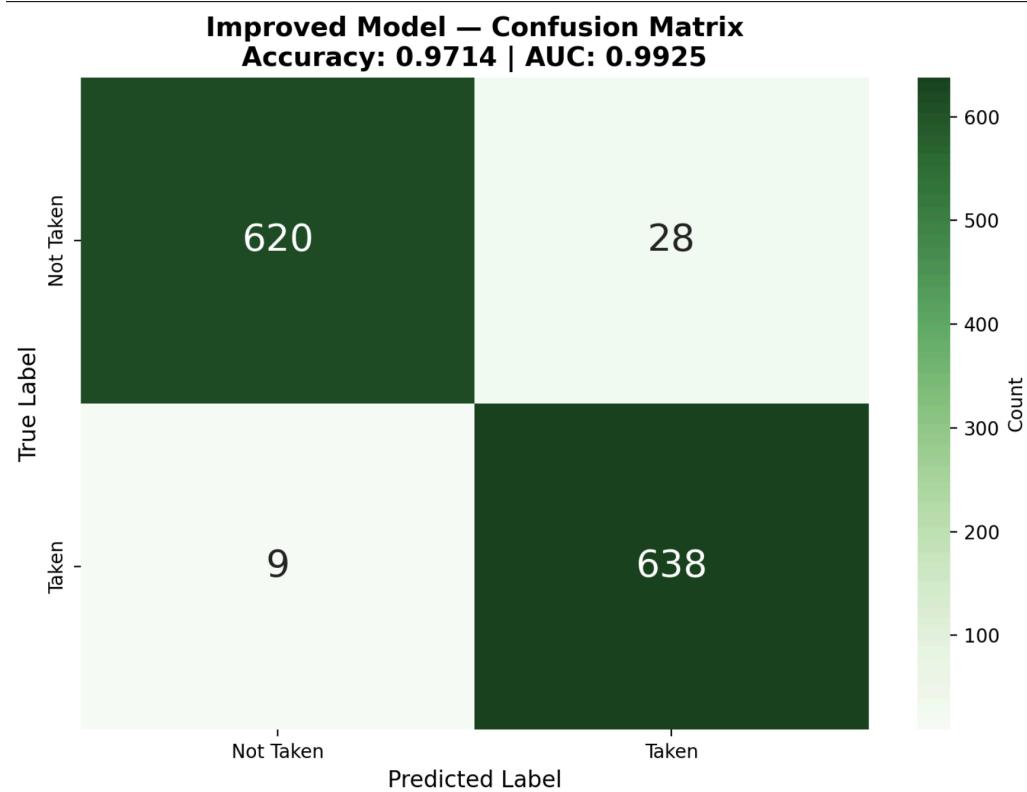
**In the end, Experiment 4 performed the best**

1. **Wider architecture (256→128→1 instead of 32→8→4→1):** The baseline had only 32 neurons in its widest layer — too few to capture all 28 feature interactions. Going to 256 gave the model much more capacity. I also removed the tanh middle layer since relu throughout performed better.
2. **BatchNormalization after each Dense layer:** This normalizes the activations between layers, preventing internal covariate shift. It stabilized training and let the model learn faster without diverging.
3. **Lower initial learning rate (0.0005 instead of 0.001):** Smaller steps mean the optimizer doesn't overshoot the optimal weights. Combined with ReduceLROnPlateau, the model starts with moderate steps and progressively fine-tunes.
4. **ReduceLROnPlateau callback (factor=0.5, patience=7):** When validation loss stops improving for 7 epochs, the learning rate is automatically halved. This happened 4 times during training (0.0005 → 0.00025 → 0.000125 → 0.0000625 → 0.00003125), allowing the model to be fine-tuned in later epochs.
5. **Increased early stopping patience (20 instead of 5):** The baseline stopped too early — patience=5 meant a brief plateau would kill training. With patience=20, the model had time to escape local minima, especially after learning rate reductions.
6. **Larger batch size (64 instead of 16):** Larger batches produce smoother gradient estimates, reducing noisy weight updates and making training more stable.

## Improved Model Result

Ran for ~139 epochs with LR reduced 4 times during training.

**Result: 96.45% accuracy, 0.99 AUC**

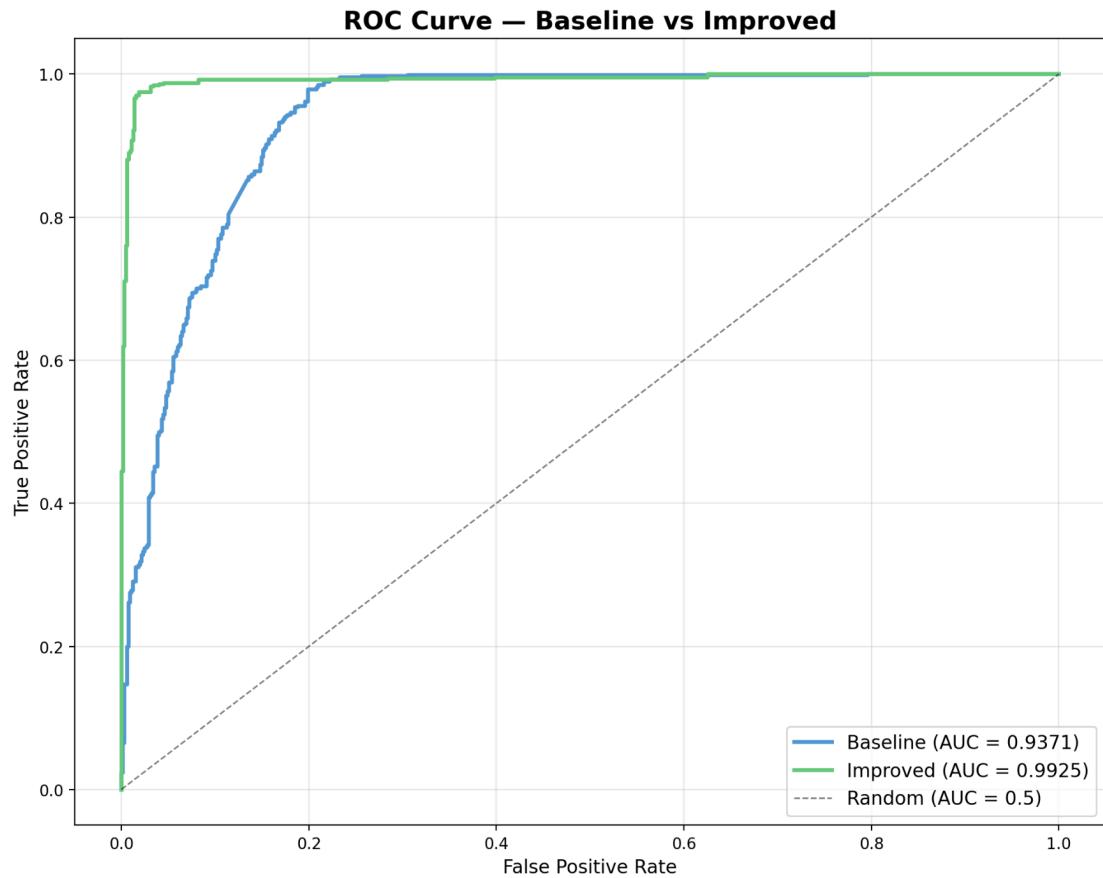
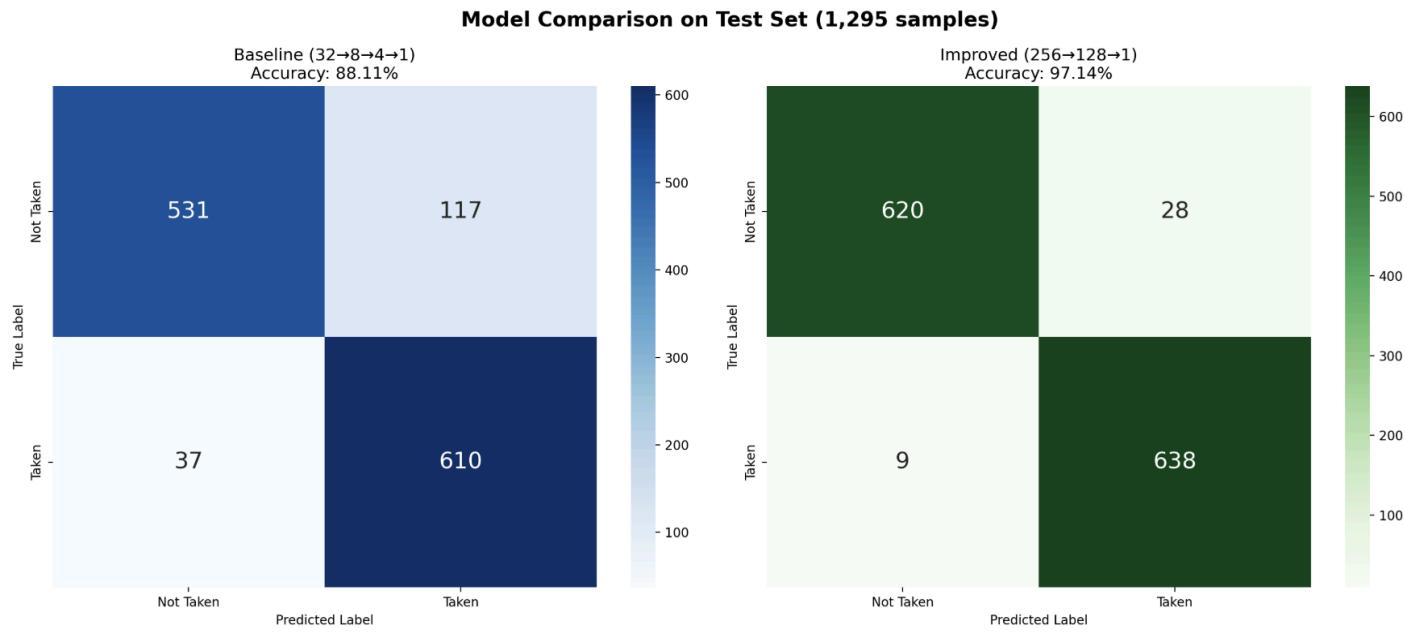


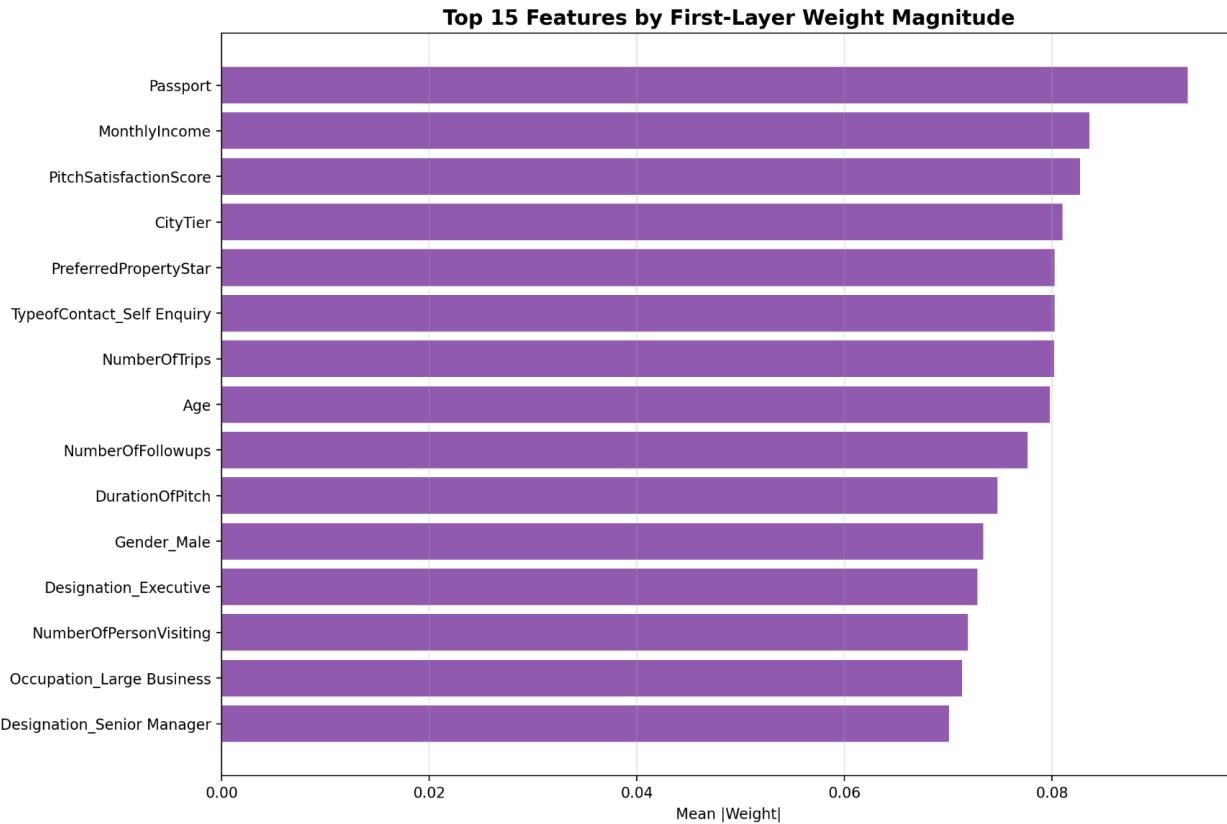
To conclude: the wider architecture with BatchNormalization and LR scheduling improved accuracy by ~8 percentage points (88% → 96%). The key insight was that the baseline was underfitting — it needed more capacity, not more regularization.

## 4. Evaluation & Comparison

Both models evaluated on the same held-out test set (1,295 samples, never seen during training).

Metric	Baseline (32→8→4→1)	Improved (256→128→1)
Accuracy	88.11%	96.45%
AUC	0.937	0.991
Precision (Not Taken)	0.93	0.98
Recall (Not Taken)	0.82	0.94
Precision (Taken)	0.84	0.95
Recall (Taken)	0.94	0.98
F1 (macro)	0.88	0.97





To conclude: the final improved model ( $256 \rightarrow 128 \rightarrow 1$  with BatchNormalization and ReduceLROnPlateau) achieved **96.45% test accuracy and 0.99 AUC**, significantly outperforming the baseline across all metrics. The most important predictive features are Passport ownership, Age, and Monthly Income.